

Continual learning accuracy vs inference compute
(20-task CIFAR-100; @128 CPG/LoRA points are bit-exact zero forgetting)

