


Continual learning accuracy vs inference compute
(20-task CIFAR-100; all @128 points are bit-exact zero forgetting)



- CViT-CPG @32 (partial pretrained transfer)
- VGG16-CPG (repro, scratch)
- CPG paper (VGG16)
- VGG16-CPG (ImageNet-pretrained)
- Linear probe @128 (frozen features + head)
- CViT-CPG @128 (full pretrained transfer, mean ± std, 3 seeds)
- Per-task LoRA r=8 @128 (frozen backbone baseline)
- Floor-only control @128 (no mechanism, 3 seeds)